

Daksh Teckchandani

608-421-4546 | dteckchandani@wisc.edu | linkedin | github

EDUCATION

University of Wisconsin–Madison

B.S. in Computer Science, Data Science, Economics

Madison, WI

Expected May 2028

- GPA: 3.67/4.0; Dean's List all semesters
- Relevant Coursework: Introduction to Algorithms, Programming III, Intro to Artificial Intelligence, Big Data Systems, Data Science Modeling I & II, Discrete Mathematics, Elementary Matrix & Linear Algebra, Intermediate Microeconomic & Macroeconomic Theory, Markets and Models, Intro to Computer Engineering

EXPERIENCE

Software Developer

CMIT NYCE

June 2025 – Present

New York City (Remote)

- **Ticket Automation:** Built custom YAML-based administration automations in PIA to auto-close minor tickets, eliminating an estimated 50–100 manual tickets/month across the client base
- **Legacy Platform Modernization – Backend & Infra:** Migrated data and integrated Microsoft Entra ID SSO for live-site user authentication (retaining a separate local master admin account for developer-level superadmin access) on a 2015-era Ruby on Rails/PostgreSQL platform; containerized with Docker and deployed via K3s on an Azure VM using ACR and a GitHub Actions CI/CD pipeline
- **Legacy Platform Modernization – Admin & UI:** Rebuilt the admin panel with flexible CRUD forms for product, pricing, and language management plus configurable multi-tier shipping logic, and redesigned the full customer-facing UI component-by-component (splash, checkout, order-confirmation flows)
- **Security Report Automation (In Progress):** Developing an end-to-end automation pipeline for client security report generation: structures raw output from an existing collection script (per-client .txt summaries and supporting .csv files, 30+ files/client across 100+ clients) and feeds it into a locally-hosted, fine-tuned LLM (Ollama) to auto-generate full 10-page client security reports, with a UI layer to review and edit the underlying scripts, targeting a reduction from the current 10–11 hours/report (2-person team) to seconds per client

PROJECTS

Statistical Strategy Validation Framework

Python, multiprocessing, Monte Carlo methods

November 2025 – Present

Personal Project

- Built a strategy library and parallelized backtesting engine with a full stress-test suite (regime-conditioned Monte Carlo simulation, regime tagging, walk-forward testing, historical risk assessment across all available years of data) covering 20 base strategies (quant, ICT, classic technical-indicator confluences), expanded via indicator-parameter permutations to 100 strategy variants, distributed across CPU cores
- Built a monitoring layer that computes and structures risk-adjusted performance metrics (Sharpe, Calmar, max drawdown, trailing 1-year risk) and a composite grader score for each variant, ranking all 100 candidates
- Built an execution layer with a metric-based selector and safelock-driven rotation system: deployed the top 10 ranked strategies (remaining 90 held as bench) into a tiered rotation, paper-traded live for 2 weeks via Sierra Chart's data feed
- Actively building a retrieval-augmented (RAG) system that ingests a raw hypothesis and a curated research library to iteratively structure and refine it into a tradable strategy specification, validating quantitative hypotheses upstream of the testing pipeline

TECHNICAL SKILLS

Languages/Libraries/Frameworks/Tools: Python, Java, JavaScript, HTML/CSS, Ruby, Pandas, NumPy, Ruby on Rails, Ollama, Git

Cloud & DevOps: Azure (VMs, ACR), Docker, Kubernetes (K3s), GitHub Actions, Microsoft Entra ID

Databases: PostgreSQL

Relevant Concepts: Statistical Modeling, Machine Learning, Monte Carlo Simulation, Data Structures & Algorithms, CI/CD